

Learning Humanoid Robot Motions through Deep Neural Networks

Luckeciano C. Melo

Autonomous Computational Systems Lab
Computer Science Division
Aeronautics Institute of Technology
São José dos Campos, Brazil
luckeciano@gmail.com

Marcos R. O. A. Maximo

Autonomous Computational Systems Lab
Computer Science Division
Aeronautics Institute of Technology
São José dos Campos, Brazil
mmaximo@ita.br

Adilson Marques da Cunha

Computer Science Division
Aeronautics Institute of Technology
São José dos Campos, Brazil
cunha@ita.br

Abstract—Controlling a high degrees of freedom humanoid robot is acknowledged as one of the hardest problems in Robotics. Due to the lack of mathematical models, an approach frequently employed is to rely on human intuition to design keyframe movements by hand, usually aided by graphical tools. In this paper, we propose a learning framework based on neural networks in order to mimic humanoid robot movements. The developed technique does not make any assumption about the underlying implementation of the movement, therefore both keyframe and model-based motions may be learned. The framework was applied in the RoboCup 3D Soccer Simulation domain and promising results were obtained using the same network architecture for several motions, even when copying motions from another teams.

Index Terms—Robotics, Machine Learning, Neural Networks.

I. INTRODUCTION

RoboCup Soccer 3D Simulation League (Soccer 3D) is a particularly interesting challenge concerning humanoid robot soccer. It consists of a simulation environment of a soccer match with two teams, each one composed by up to 11 simulated NAO robots [8], the official robot used for RoboCup Standard Platform League since 2008. Soccer 3D is interesting for robotics research since it involves high level multi-agent cooperative decision making while providing a physically realistic environment which requires control and signal processing techniques for robust low level skills.

In the current level of evolution of Soccer 3D, motion control is a key factor in team's performance. Indeed, controlling a high degrees of freedom humanoid robot is acknowledged as one of the hardest problems in Robotics. Much effort has been devised to humanoid robot walking, where researchers have been very successful in designing control algorithms which reason about reduced order mathematical models based on the Zero Moment Point (ZMP) concept, such as the linear inverted pendulum model [9]. Nevertheless, these techniques restrict the robot to operate under a small region of its dynamics, where the assumptions of the simplified models are still valid [4], [15].

Therefore, model-based techniques are hard to use for designing highly dynamic movements, such as a long distance kick and a goalkeeper's dive to defend the goal from a fast moving ball. In the robot soccer domain, a common approach

for these movements is to employ keyframe movements, where the motion is composed by a sequence of robot postures. In this case, the movement is designed off-line and executed in a open-loop fashion in execution time.

Due to the lack of mathematical models, an approach frequently employed is to rely on human intuition to design keyframe movements by hand, usually aided by graphical tools. However, this process is difficult, time consuming, and is often unable to obtain high performance motions given the high dimensionality of the search space. Other possible solution is to use motion capture data from humans [18], which has its own challenges due to the fact that the kinematic and dynamic properties of a humanoid robot differs greatly from those of a human.

Therefore, many works have experimented on using machine learning and optimization algorithms to develop high performance keyframe movements, showing promising results. Rei et al. describes an algorithm which is able to mimic movements observed from other agents and improve these learnt motions through an evolutionary strategy [5]. Other works have also used optimization to improve the performance of existing keyframe movements [1], [15]. Moreover, Peng et al. have developed control policies through deep reinforcement learning that mimic reference motions [16].

A common fact regarding the aforementioned machine learning and optimization approaches is that they rely on reference motions for learning. Due to the high dimensionality of the search space, direct optimization from a completely random point most likely fails to find an useful movement. In this work, we contribute by showing that a neural network may be taught through supervised learning to mimic an existing movement. From an engineering standpoint, having a keyframe movement represented as a neural network does not provide advantages by itself. However, our intention in a future work is to use this neural network as a seed for reinforcement learning methods as the one shown in [16].

The remaining of this work is organized as follows. Section II provides theoretical background. In Section III, the methodology and tools used in this work are explained. Furthermore, Section IV presents simulation results to validate our approach. Finally, Section V concludes and shares our ideas for future

work.

II. BACKGROUND

A. Keyframe Movements

Definition 1: A *keyframe* $\mathbf{k} = [j_1, j_2, \dots, j_n]^T \in K \subseteq \mathbb{R}^n$ is an ordered set of joint angular positions, where K and n are the joint space and the number of degrees of freedom of the robot, respectively.

Definition 2: A *keyframe step* is an ordered pair $\mathbf{s} = (\mathbf{k}, t) \in S = K \times \mathbb{R}$, where \mathbf{k} is a keyframe and t is the time when the keyframe must be achieved with respect to the beginning of the movement, respectively.

Definition 3: A *keyframe movement*, or simply a movement, is defined as $\mathbf{m} = (\mathbf{s}_1, \mathbf{s}_2, \dots, \mathbf{s}_\gamma, r) \in M = S^\gamma \times \mathbb{R}$, where γ and r are the number of keyframe steps and the speed rate of the movement, respectively. In this representation, we assume the movement starts at time 0 and the first keyframe step represents the robot posture at the beginning of the movement. Therefore, $t_1 = 0$ and each time $t_i, \forall i \geq 2$ is a time since the beginning of the movement.

Keyframe movements are executed in an open-loop fashion, where joint positions are computed through interpolation of keyframe steps based on the current time. If the interface to the robot joints is not position-based, local controllers may be used to track the position references issued by the keyframe. For example, in the Simspark simulator, the simulated NAO has speed-controlled joints, therefore we use simple proportional controllers for each joint to track the desired joint positions. To obtain smooth joint trajectories, we interpolate keyframe steps using cubic splines [2], which are functions of class \mathcal{C}^2 .

B. Neural Networks

Neural Networks are a learning representation whose goal is to approximate some function f^* . The data collected from an environment encodes an underlying function $\mathbf{y} = f^*(\mathbf{x})$ that maps an input \mathbf{x} to an output \mathbf{y} , which may be a category from a classifier or a continue value in regression problems. The neural network defines an approximate mapping $\mathbf{y} = f(\mathbf{x}; \boldsymbol{\theta})$ by learning the values of the parameters $\boldsymbol{\theta}$ which result in the best function approximation. The Figure 1 shows a neural network and an artificial neuron in detail.




Fig. 1. Artificial neuron and feed forward artificial neural network [19].

These networks are typically represented by composing together many different functions, which are associated with a directed acyclic graph describing a computational model. For example, we might have three layers (each of them representing a function $f^{(1)}, f^{(2)}$, and $f^{(3)}$ respectively), connecting in a chain, resulting in a final representation $f(\mathbf{x}) = f^{(3)}(f^{(2)}(f^{(1)}(\mathbf{x})))$.

During neural network training, the objective is to adjust $f(\mathbf{x})$ to match $f^*(\mathbf{x})$ using the training dataset, which provides noisy examples of $f^*(\mathbf{x})$ evaluated in different points. The training examples specify directly what the output layer must do at each point \mathbf{x} , but the learning algorithm must decide how to use all layers to produce this desired output [7].

Additionally, we must also choose a learning algorithm to tune this function approximation. In the context of neural networks, gradient-based algorithms are broadly used, especially those based on the backpropagation idea [17]. The purpose of these algorithms are to propagate the gradient of a cost function through the whole network, in order to minimize the cost function. Most modern neural networks perform this optimization strategy using maximum likelihood, i.e. the cross-entropy between the training data and the model distribution:

$$J(\boldsymbol{\theta}) = -\mathbb{E}_{\mathbf{x}, \mathbf{y} \sim \hat{p}_{data}} \log p_{model}(\mathbf{y}|\mathbf{x}) \quad (1)$$

In this work, we used the mean squared error loss function in order to fit the dataset. Indeed, we may show that both cost functions are closely related. Let us consider normally distributed errors:

$$p_{model}(\mathbf{y}|\mathbf{x}) = \mathcal{N}(\mathbf{y}; f(\mathbf{x}; \boldsymbol{\theta}), \sigma^2 \mathbf{I}) \quad (2)$$

where $f(\mathbf{x}; \boldsymbol{\theta})$ and $\sigma^2 \mathbf{I}$ are the mean and covariance of this distribution, respectively. Substituting Eq. (2) in Eq. (1):

$$J(\boldsymbol{\theta}) = \frac{1}{2} \mathbb{E}_{\mathbf{x}, \mathbf{y} \sim \hat{p}_{data}} \|\mathbf{y} - f(\mathbf{x}; \boldsymbol{\theta})\|^2 + const \quad (3)$$

The constant term does not depend on $\boldsymbol{\theta}$ and may be dropped. By explicitly evaluating the expectation in Eq. (3), we arrive at the mean squared error cost function:

$$J(\boldsymbol{\theta}) = \frac{1}{2m} \sum_i^m \|y_i - f(\mathbf{x}; \boldsymbol{\theta})\|^2 \quad (4)$$

Lastly, the gradient of the loss function is taken and propagated through the hidden layers by the chain rule. For example, given $\mathbf{Y} = g(\mathbf{X})$ and $z = f(\mathbf{Y})$, then the chain rule states:

$$\nabla_{\mathbf{x}} z = \sum_j (\nabla_{\mathbf{x}} Y_j) \frac{\partial z}{\partial Y_j} \quad (5)$$

This equation is taken recursively until the gradient is propagated to all layers of the neural network.

III. METHODOLOGY AND TOOLS

A. Dataset

In order to use supervised learning for learning keyframe motions using neural networks, we first need to construct a dataset. A dataset consists of samples of keyframe steps. The samples were collected within Soccer 3D environment with a frequency of 50 Hz. We acquired these samples in two different ways.

In the first one, we commanded an agent of our team to execute specific motions and sampled the reference joint positions computed by our code. In this case, we sampled the kick and get up keyframe motions [15]. Notice that for this approach to be successful, one needs access to the source code.

The second approach involved changing the Soccer 3D server source code to provide current joint positions of a given robot, in a similar way as described in [12]. This allowed us to acquire motion datasets from other teams, without any knowledge of how these movements are implemented. In this case, we collected two types of kicks based on keyframes and sampled joint values of the walking engine [11].

B. Neural Network Architecture and Hyperparameters

The neural network has to be able to learn how to interpolate between samples, which actually happens. The architecture that performed best – in terms of mean absolute error minimization and simplicity – is shown in Figure 2. A deep neural network with 2 hidden, fully connected layers of 75 and 50 neurons was used. The output layer has 23 regression neurons, which represent the 22 joint angles and a neuron whose output indicates if the motion has ended or not. The neurons in each hidden layer use the LeakyReLU activation function [20]:

$$f(x) = \begin{cases} \alpha x, & x < 0 \\ x, & x \geq 0 \end{cases}$$

where α is a small constant. This activation function was used to improve the representation capacity of the neural network, adding support for non-linear functions.




Fig. 2. Architecture of the neural network designed to learn motions.

This architecture resulted in thousand of parameters to optimize, as exposed in Table I. A very high number when compared to more traditional optimization approaches [11].

Notice that increasing the number of parameters usually allows representing better movements.

TABLE I
NETWORK SUMMARY

Layer	Neurons	Activation	Parameters
Dense	75	LeakyReLU	150
Dense	50	LeakyReLU	3800
Dense	23	Linear	1173
Total Parameters			5123

C. Training Procedure

Since keyframe motions are executed in an open-loop fashion, the sequence of joint positions are always the same for different repetitions, independently of the robot's state. Therefore, adding samples of multiple executions of the same motion would not make our dataset richer, so we decided to use only one repetition for each movement for faster training. In the case of the walking motion, we collected samples within one walking period.

During training, we used 50 thousands epochs divided in 5 training phases, where the learning rate was decreased between phases in order to achieve better performance. First, we executed 30000 epochs using learning rate of 0.001. The other phases had 5000 epochs each, and we decreased the learning rate by 0.0002 in each phase.

Furthermore, we used Adam optimization [10] during the whole training. The loss function used was the mean squared error, as explained in Subsec. II-B. We decided this loss function is adequate for this problem because it strongly penalizes large errors, which can collapse the whole motion.

D. Deployment in Soccer 3D Environment

In order to perform network design and the training procedure, we used the Keras [3] framework coupled with Tensorflow [6] as backend. After training, the weights were freedzed and converted to a specific format which is readable using the Tensorflow C++ API integrated within the agent's code. Hence, the training is performed outside the environment, but the agent actually computes network inferences during simulation execution.

IV. RESULTS AND DISCUSSION

A. Training Results

All results and logs obtained, as well as the code used for neural network training, are available in the project repository¹ for reproducibility.

The initial results come from the training procedure, outside the simulation environment. Figure 3 presents training curves for the kick keyframe dataset. In this case, the plots show mean squared error and mean absolute error metrics, respectively. In both metrics, the value decreases drastically in the first epochs. This same behavior was present in other training procedures as well. However, only after thousands of epochs the network

¹ Repository: <https://goo.gl/MjRWAH>




Fig. 3. The plots of mean squared error and mean absolute error during training.

achieved a low error that reproduced the motion successfully, which shows how sensible to small joint errors keyframes are, given that they are open-loop motions. The peaks during the training happen at the learning rate transition instants, but they do not hurt the training procedure itself. This is due to the fact that we re-compile the model at each training phase, which resets the optimizer state. This means that the training will suffer a little at the beginning until you adjust the learning rate, the moments, etc. However, there is absolutely no damage to the weights.

B. The Learned Kick Motion

The final mean absolute error is **0.018** radians and the motion is visually indistinguishable from the original one, as can be seen in Figure 4. In this figure, snapshots from both motions were taken. The Figure 5 shows several plots of joint angles comparing the original and learned kick motions. As we may see, the learned motion has fitted the movement with minor errors².




Fig. 4. Kick motion. The first row of figures shows the original kick motion. The second row shows the learned kick motion. The motions are visually indistinguishable.

In order to evaluate the learned kick motion in the RoboCup Soccer 3D domain, we created a statistical test. Inside the test scenario, the ball was placed initially in the center of the field with an agent near to it. The only action of the agent is to kick

the ball in the goal direction. After the kick, the agent runs until reaching the ball and kicks it again, repeating this process till scoring a goal. When the goal occurs, this same scenario is repeated. The whole test was conducted during thirty minutes in clock time and the following data was collected: total number of kicks, number of successful kicks, mean distance that the ball has traveled and the standard deviation of this measure. The results from the original and learned kicks is shown in Table II.

TABLE II
KICK COMPARISON

Kick Type	Statistics		
	Accuracy (%)	Distance (m)	
		Mean	Std
Original Kick	64.5	8.92	3.82
Neural Kick	52.6	7.16	4.06

Although both kicks have similar results, the original kick is slightly better in this scenario. Confronting Figure 5, we can conclude that even with an almost equal representation, the kick lose part of its efficiency and this fact show us how sensible are movements based on keyframe data.

C. The Learned Walk Motion

Using the modified server described in Subsec. III-A, a dataset with samples of the UT Austin Villa's walking motion [12] was acquired. This team is the current champion of RoboCup Soccer 3D competition [14].

The objective is to mimic the walk motion as a keyframe and use that in our agent. The previously described framework used for learning our own kick motion was used in this training, including the neural network architecture and its hyperparameters.

The results from this training are shown in Figure 6. Similarly to Figure 5, it shows the joint angles throughout the walking motion period for the original and learned walk. Additionally, it shows the real joints values from the movement in the server. These joints were chosen because they are the most dynamic in the walk motion and therefore the hardest to learn.

The learned motion has fitted the dataset very well. However, these values are just desired joints. In fact, these values are used as references to joint controllers and are attenuated due to joint dynamics. Furthermore, this motion is operated in a open-loop fashion, so the agent is not able to correct its own trajectory, and this walks gets biased in the simple task of walking straight forward.

Despite the facts above, the motion works well in a non-competitive scenario³, which is shown in the metrics collected in Forward Walk test scenario – agent walking forward from the goal post until the center line of field – in Table III and the visual representation in Figure 7.

² Kick results video: <https://streamable.com/gpltm>

³ Walk results video: <https://streamable.com/m5w1d>

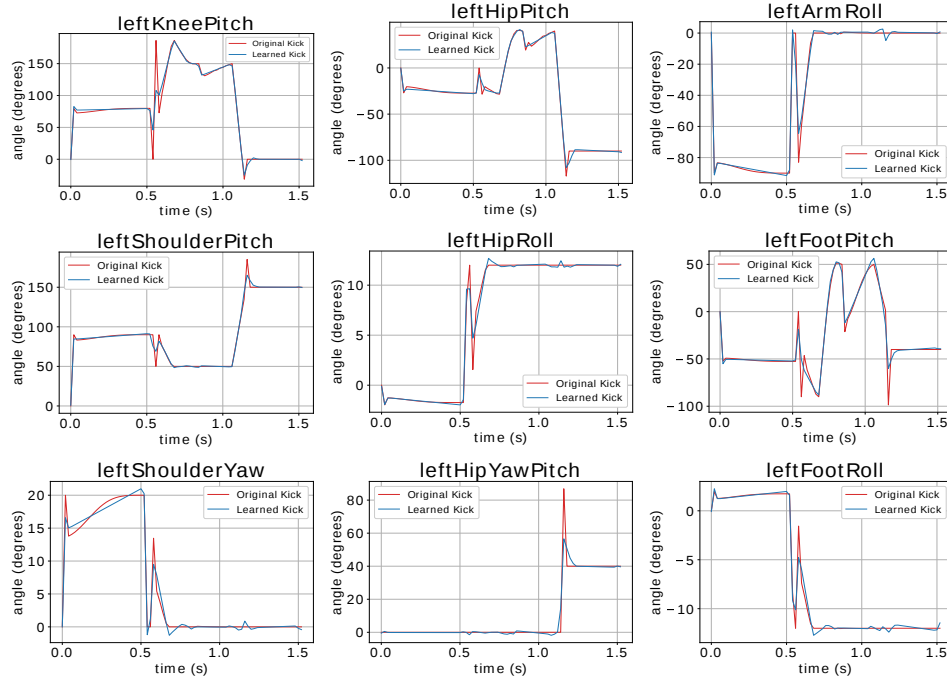


Fig. 5. Joint values for comparing original and learned kicks. The neural network was able to fit the joint trajectories with small errors.

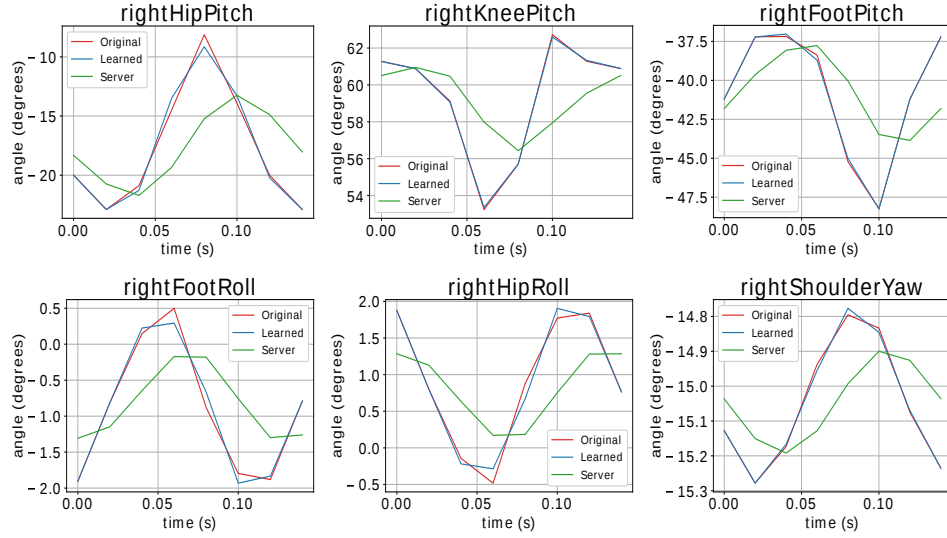


Fig. 6. Joints positions during a period of the walking motion for the original walk, the learned walk and the joints positions effectively attained during the learned walking motion.

TABLE III
WALK COMPARISON - FORWARD WALK

Walk Type	Statistics			
	Velocity (m/s)		Y Error (m)	
	Mean	Std	Mean	Std
Original Walk	0.87	0.01	-	-
Learned Walk	0.23	0.01	0.96	2.63

D. Other motions

This same framework was used to learn other keyframe motions originated by our agent itself, such as the get up

motion. As the cases previously described, the resultant neural network was capable of mimicking the keyframe including its interpolation. Hence, all of our keyframe motions can be replaced by neural motions with similar performance.

However, the huge improvement of this method is about mimicking other teams motions. In Soccer 3D environment, movements like kick and walking have giant impact in the team's performance. With this learning framework, our agent is able to mimic multiples movements from several teams.

As example, we collected data from UT Austin Villa kick,




Fig. 7. Walking motions comparison. The figure (a) shows our agent in its regular walk. The figure (b) shows the same agent mimicking UT Austin Villa walk. The figure (c) is the UT Austin Villa agent itself, performing his own walking motion.

which was originally optimized using Deep Reinforcement Learning techniques [13]. Our agent learned this kick without any additional optimization strategy: we just used samples collected from the modified server.

V. CONCLUSION AND FUTURE WORK

In this work, we presented a method for learning humanoid robot movements using datasets composed of joint values at each time instant. The learning framework provided was capable to learn several types of motion, including walk and kick, without any change in network architecture or hyperparameters. Moreover, the learned motions have similar performance to the original ones. Furthermore, this framework was able to learn other teams motions, without any knowledge about the underlying implementation – only using the joints values provided by a modified version of the server. This is a huge improvement in terms of getting improved motions, as our agent can mimic other teams motions using this machine learning technique.

As future work, we plan to apply Deep Reinforcement Learning algorithms to obtain faster and more robust kicks, using as “seed” the neural networks obtained in this work. Another track to be followed is to transfer the learning of this network to a new network that represents the motion policy itself (i.e a network which has as inputs the current state of the robot, including joint and link states, besides the current time instant), and optimize this motion policy in order to get a closed-loop walking and kick motion that can correct itself. As a long term goal, we intend to create model-free kick and walking engines.

VI. ACKNOWLEDGEMENTS

We thank our sponsors ITAEx, Altium, Intel, Mathworks, Metinjo, Micropress, Poliedro, Polimold, Pouplex, FHC, Rapid, and Solidworks. We are also grateful to ITA for supporting our work.

We would also like to show our gratitude to Patrick MacAlpine from UT Austin Villa team for sharing their ideas and code regarding the Soccer 3D simulation server modification.

Finally, we would like to acknowledge all the ITAndroids team, especially Soccer 3D simulation team members for the hard work in the development of the base code.

REFERENCES

- [1] Abbas Abdolmaleki, David Simões, Nuno Lau, Luis Paulo Reis, and Gerhard Neumann. Learning a humanoid kick with controlled distance. In Sven Behnke, Raymond Sheh, Sanem Sarel, and Daniel D. Lee, editors, *RoboCup 2016: Robot World Cup XX*, pages 45–57, Cham, 2017. Springer International Publishing.
- [2] Richard H. Bartels, John C. Beatty, and Brian A. Barsky. *An Introduction to Splines for Use in Computer Graphics & Geometric Modeling*. Morgan Kaufmann Publishers Inc., San Francisco, CA, USA, 1987.
- [3] François Chollet et al. Keras. <https://keras.io>, 2015.
- [4] Steven Collins, Andy Ruina, Russ Tedrake, and Martijn Wisse. Efficient Bipedal Robots Based on Passive Dynamic Walkers. *Science Magazine*, 307:1082–1085, February 2005.
- [5] Mike Depinet, Patrick MacAlpine, and Peter Stone. Keyframe sampling, optimization, and behavior integration: Towards long-distance kicking in the robocup 3d simulation league. In H. Levent Akin, Reinaldo A. C. Bianchi, Subramanian Ramamoorthy, and Komei Sugiura, editors, *RoboCup-2014: Robot Soccer World Cup XVIII*. Springer Verlag, 2015.
- [6] Martín Abadi et al. TensorFlow: Large-scale machine learning on heterogeneous systems, 2015. Software available from tensorflow.org.
- [7] Ian Goodfellow, Yoshua Bengio, and Aaron Courville. *Deep Learning*. MIT Press, 2016.
- [8] D. Gouaillier, V. Hugel, P. Blazevic, C. Kilner, J. Monceaux, P. Lafourcade, B. Marnier, J. Serre, and B. Maisonnier. Mechatronic design of nao humanoid. In *2009 IEEE International Conference on Robotics and Automation*, pages 769–774, May 2009.
- [9] Shuuji Kajita, Fumio Kanehiro, Kenji Kaneko, Kazuhito Yokoi, and Hirohisa Hirukawa. The 3D Linear Inverted Pendulum Mode: A simple modeling for a biped walking pattern generation. In *In Proceedings of the 2001 IEEE/RSJ International Conference on Intelligent Robots and Systems*, Hawaii, USA, October 2001. IEEE.
- [10] Diederik P. Kingma and Jimmy Ba. Computer science & learning adam: A method for stochastic optimization. *arXiv*, Dec 2014.
- [11] Patrick MacAlpine, Samuel Barrett, Daniel Urieli, Victor Vu, and Peter Stone. Design and optimization of an omnidirectional humanoid walk: A winning approach at the RoboCup 2011 3D simulation competition. In *Proceedings of the Twenty-Sixth AAAI Conference on Artificial Intelligence (AAAI)*, July 2012.
- [12] Patrick MacAlpine, Nick Collins, Adrian Lopez-Mobilia, and Peter Stone. Ut austin villa: Robocup 2012 3d simulation league champion. In Xiaoping Chen, Peter Stone, Luis Enrique Sucar, and Tijn van der Zant, editors, *RoboCup 2012: Robot Soccer World Cup XVI*, pages 77–88, Berlin, Heidelberg, 2013. Springer Berlin Heidelberg.
- [13] Patrick MacAlpine and Peter Stone. Ut austin villa 3d simulation soccer team 2017. 2017.
- [14] Patrick MacAlpine and Peter Stone. UT Austin Villa: RoboCup 2017 3D simulation league competition and technical challenges champions. In Claude Sammut, Oliver Obst, Flavio Tonidandel, and Hidehisa Akyama, editors, *RoboCup 2017: Robot Soccer World Cup XXI*, Lecture Notes in Artificial Intelligence. Springer, 2018.
- [15] F. Muniz, M. R. O. A. Maximo, and C. H. C. Ribeiro. Keyframe movement optimization for simulated humanoid robot using a parallel optimization framework. In *2016 XIII Latin American Robotics Symposium and IV Brazilian Robotics Symposium (LARS/SBR)*, pages 79–84, Oct 2016.
- [16] Xue Bin Peng, Pieter Abbeel, Sergey Levine, and Michiel van de Panne. Deepmimic: Example-guided deep reinforcement learning of physics-based character skills. *ACM Transactions on Graphics (Proc. SIGGRAPH 2018 - to appear)*, 37(4), 2018.
- [17] David E. Rumelhart, Geoffrey E. Hinton, and Ronald J. Williams. Learning representations by back-propagating errors. *Nature*, 323(533), 1988.
- [18] Aaron P. Shon, Keith Grochow, and Rajesh P. N. Rao. Robotic imitation from human motion capture using gaussian processes. In *In Proceedings of the IEEE/RAS International Conference on Humanoid Robots (Humanoids)*, 2005.
- [19] Dejan Tanikic and Vladimir Despotovic. Artificial intelligence techniques for modelling of temperature in the metal cutting process. 2012.
- [20] Bing Xu, Naiyan Wang, Tianqi Chen, and Mu Li. Empirical evaluation of rectified activations in convolutional network. *arXiv*, Nov 2015.